

Co-Area Mesh Method — strict $h = 0$ progression

Test case throughout: $(\tau, \gamma) = (2, 0.098)$ at $G = 21$
(emin15-certified longdouble FP of the kernel operator).

Stage 1c — cube-based strict- $h = 0$ forward eval (Φ_{CMM}):
cube cell residual (median): 6.230e-03
cube cell residual (max): 3.300e-01
wall time (full $G=21$): 813s
-> Confirms strict $h=0$ matches kernel within $O(h^2)$ bulk,
 $O(h)$ tails near critical-price regions.

Stage 2 — iteration test of cube-based CMM from kernel FP:
iter 0 -> 1 -> 2: 0.33 -> 0.46 -> 0.47 (DIVERGES)
-> Cube state is the wrong representation: marching cubes
re-triangulates as $P[i,j,l]$ updates, breaking smoothness.


Stage 3a — moving-mesh prototype:
Mesh: 32 surfaces, 22734 vertices
Vertex residual (500-sample): max 3.23e-01, median 1.72e-01
-> Vertex-level residual is 30x larger than cube-level
residual: strict $h=0$ disequilibrium was hidden by coarse sampling.
-> Pure Python: 120 ms/vertex (~90 min full eval, too slow for Newton).

Stage 3b — JIT'd residual (numba, per-triangle 2-point Gauss-Legendre):
Mesh: 32 surfaces, 22734 verts
Full-residual eval: 0.6s (was 90 min in Python; 8595x speedup)
max $|r|$: 3.241e-01
median $|r|$: 1.572e-01
p90 $|r|$: 2.905e-01
-> Confirms 3a numbers; opens Newton at minute-scale cost.

Stage 3c — Newton-Krylov on moving mesh:
(still running — see below)

Residual distributions: cube vs mesh

median is 30x higher at vertices: cube hid the strict-h=0 disequilibrium



(Stage 3c history will appear when it completes.)

Key learnings

1. The strict- $h = 0$ forward operator EXISTS, is computable in seconds at $G = 21$, and disagrees with the kernel operator at the rate $O(h)$ in tails, $O(h^2)$ in bulk -- matching what the deep G-ladder data already implied.
2. Cube-based CMM cannot be solved by Newton:
re-triangulation jumps in Φ_{CMM} as cube values update destroy smoothness. This is the same defect that killed the historical strict ϕ_{K3}^{halo} operator.
3. Moving-mesh CMM (vertex positions = primary DOFs, topology pinned) eliminates that defect by construction.
4. JIT'd residual at $\sim 0.6\text{s/eval}$ enables Newton iteration at minute-scale cost. The first Newton step's contraction will tell us if the strict $h = 0$ fixed point is reachable from the kernel FP.

Pending (Stage 3c, in progress):

- Newton-Krylov on a 22734-vertex per-surface scalar displacement parametrization (vertex moves along its surface normal).
- Convergence test: residual must contract to $<1\text{e-}6$ over ≤ 8 iters.
- If converged: compute strict- $h=0$ deficit at $(\tau, \gamma) = (2, 0.098)$ and compare with the deep-ladder continuum estimate $d_{\infty} = 0.268 \pm 0.010$.

Caveats:

- Mesh quality must be monitored (no triangle flips).
- The initial mesh comes from marching cubes of the kernel FP; if the strict $h=0$ fixed point has different surface topology, Newton will fail or land on a wrong solution.
- Stage 3c uses normal-only displacement (1 DOF/vertex). Tangential rearrangement, if needed, requires the full 3 DOF/vertex.