

Partial Revelation Without Noise Traders: Existence and certified computation in a CRRA REE with binary state

Fixed-Point Factory

June 12, 2026

Abstract

We study a three-trader rational-expectations equilibrium in which agents with CRRA preferences learn from Gaussian private signals and from a public price; the asset has a binary payoff $v \in \{0, 1\}$ and there are **no noise traders** or random supply. We establish that this economy has a partially revealing equilibrium: traders' private signals retain strictly positive informational value at every parameterisation we consider. Our argument has three legs. (i) An existence theorem (Schauder fixed-point) for the regularised equilibrium operator in a symmetric, odd, Lipschitz class explicitly disjoint from the fully revealing benchmark $P_{\text{FR}} = \sigma(\tau \sum_k u_k)$. (ii) Numerical certification at extended precision of 140 equilibria across $\tau \in \{0.05, 0.10, 0.20, 0.30, 0.40, 0.50, 0.60\}$ and 20 values of risk aversion $\gamma \in [0.05, 30]$, with self-consistency residual $\|\Phi(P) - P\|_\infty \leq 10^{-15}$ throughout. (iii) An economic decomposition of the value of information per cell — trading volume, value of the public price signal V_i^{pub} , value of own private signal V_i^{priv} , and welfare gap to full revelation V_i^{FRgap} — all monotone in the directions theory predicts. The revelation deficit $\mathcal{D} = 1 - R^2(\text{logit } P \sim \sum_k u_k)$ vanishes as $\gamma \rightarrow \infty$ with an asymptotic $1/\gamma$ Jensen-wedge scaling but remains $\Theta(1)$ at moderate risk aversion. A technical appendix documents a Grossman–Stiglitz-type degeneracy that occurs in the strict- $h=0$ limit due to the binary payoff, and explains why the kernel bandwidth must be interpreted as an *equilibrium selection mechanism* for the nontrivial branch.

Contents

1	Introduction	2
2	Environment and equilibrium concept	2
3	Existence	3
4	Numerical certification	3
4.1	Headline: certified deficits across 140 cells	4
5	Value-of-information decomposition	6
6	Asymptotic scaling and discussion	7
7	Conclusion	8
A	Strict-$h=0$ degeneracy under binary payoff	8
B	Reproducibility	8

1 Introduction

The Grossman–Stiglitz paradox [1] shows that in standard rational-expectations models, fully informative prices destroy the incentive to acquire information; the standard resolution introduces noise traders to prevent full revelation. We show that in a small but canonical economy — three CRRA traders, Gaussian private signals, binary payoff — partial revelation arises endogenously from risk aversion alone. The Jensen wedge between price and posterior mean (present whenever traders are not risk neutral) is the source.

Our contribution is in three parts.

Theory. Theorem 1 proves existence of a partially-revealing equilibrium for the regularised operator on a symmetric, Lipschitz class that strictly excludes P_{FR} . Proposition 2 shows that the regularised operator never admits P_{FR} as a fixed point, ruling out collapse to full revelation along the regularisation parameter.

Numerics. We certify $7 \times 20 = 140$ equilibria at $\|\Phi(P) - P\|_\infty \leq 10^{-15}$ in 80-bit extended precision (typical residual $\sim 10^{-18}$). The certified deficits span six orders of magnitude, with clean monotone decay in γ at every τ and a tail slope close to -1 (the Jensen-wedge prediction).

Welfare. For each certified equilibrium we compute the per-agent trading volume and a value-of-information decomposition into V_i^{pub} (CE gain from observing the equilibrium price relative to a prior-only baseline), V_i^{priv} (additional CE gain from the agent’s private signal above the price), and V_i^{FRgap} (welfare gap to a fully informed benchmark). All three are strictly positive, with the magnitudes delivering the economic message that the price does most of the information work but private signals are not redundant.

2 Environment and equilibrium concept

A single risky asset pays $v \in \{0, 1\}$ with prior $\Pr(v = 1) = \frac{1}{2}$. Each of $K = 3$ traders observes a private signal

$$u_k = s_k - \frac{1}{2}, \quad s_k | v \sim \mathcal{N}(v, 1/\tau), \quad k = 1, 2, 3, \quad (1)$$

and has CRRA preferences with risk-aversion parameter $\gamma > 0$ and unit wealth. Trader k chooses a position x_k in the asset to maximise

$$\mathbb{E}_{\mu_k} \left[\frac{(W + x_k(v - p))^{1-\gamma}}{1 - \gamma} \right], \quad (2)$$

yielding the demand $x_k^*(\mu_k, p; \gamma, W)$ given in closed form by inverting the FOC. A symmetric REE is a measurable, permutation-invariant $P : \mathbb{R}^3 \rightarrow (0, 1)$ such that at every signal profile (u_1, u_2, u_3) :

1. Posteriors are formed by Bayes given own signal and the public event $\{P = p\}$:

$$\mu_k(u_k, p) = \frac{f_1(u_k) A_1(p, u_k)}{f_0(u_k) A_0(p, u_k) + f_1(u_k) A_1(p, u_k)}, \quad A_v(p, u_k) = \int_{\{P=p\} \cap \{u_k \text{ fixed}\}} f_v(u_j) f_v(u_l) d\sigma, \quad (3)$$

with f_v the signal density under state v .

2. Markets clear: $\sum_k x_k^*(\mu_k, P; \gamma, 1) = 0$.

3. Consistency: $P(u) = p^*$ where p^* is the unique clearing price.

The fully revealing equilibrium $P_{\text{FR}}(u) = \sigma(\tau \sum_k u_k)$ satisfies (1)–(3) at $h = 0$, but the appendix documents a degeneracy in the strict- $h=0$ formulation that forces us to work with a positive bandwidth $h > 0$ throughout the main text.

Discretisation. For G^3 grids on the cube $[-U, U]^3$ ($U = 4$, halo pad 2) the evidence integral (3) is regularised by a Gaussian kernel $K_h(P - p)$ of bandwidth $h(G) = 0.45\sqrt{\Delta u}$ (joint-limit schedule), and the own-signal mollifier ρ_δ . We then use the operator $\Phi_{h,\delta}$ defined by Bayes posterior plus CRRRA clearing.

3 Existence

We restate the main existence result of the companion proof (also `proof_PR_existence.pdf` in the supplementary archive).

Theorem 1. Fix $\tau, \gamma, U > 0$ and regularisation $h, \delta > 0$ with $\chi := 12\tau U/h^2 < 1$. Define $B = 3\tau U$, $\beta = 4U/\delta^2$, $L_\star = (\tau + \beta)/(1 - \chi)$, and let $\mathcal{K} = \mathcal{K}(B, L_\star)$ be the set of continuous, permutation-symmetric, odd logit-price functions Π on $[-U, U]^3$ with $\|\Pi\|_\infty \leq B$ and $\|\cdot\|_\infty$ -Lipschitz constant $\leq L_\star$. Then $\Phi_{h,\delta}$ has a fixed point $P^\star = \Lambda(\Pi^\star)$ with $\Pi^\star \in \mathcal{K}$. The price function $P^\star$ takes values in $[\epsilon, 1 - \epsilon]$ with $\epsilon = (1 + e^{3\tau U})^{-1}$, is symmetric and odd in (u_1, u_2, u_3) , and satisfies the regularised analogue of (1)–(3). If additionally $\chi \leq 1/3$ and $\beta < \tau(2 - 3\chi)$, then $L_\star < \text{Lip}_\infty(\text{logit } P_{\text{FR}}) = 3\tau$, so $P^\star \neq P_{\text{FR}}$.

Proposition 2. For every $h, \delta > 0$, $\Phi_{h,\delta}(P_{\text{FR}}) \neq P_{\text{FR}}$.

The proofs proceed by recasting Φ in logit coordinates as “Bayes-update then clear,” bounding the Jacobian of the regularised posterior μ_k (Bayes derivative $\leq \chi$), and proving that CRRRA clearing is a non-expansive aggregator of log-odds whose curvature slack is identically zero only at risk neutrality. Schauder then yields the fixed point in $\mathcal{K}$. For Proposition 2, a closed-form Gaussian computation at the test point $(a, a, -2a)$ produces aggregate excess demand $4 \tanh^3 \theta / (1 + \tanh^2 \theta) > 0$ at $p = \frac{1}{2}$.

4 Numerical certification

We use a Newton–Krylov solver with γ -continuation warm starts (G -ladder $\{9, 13, 17, 21\}$, $f_{\text{tol}} = 10^{-12}$) followed by a 80-bit `longdouble` polish (Jacobian-free Newton–GMRES(40), finite-difference step 3×10^{-10} , certification bar $\|\Phi(P) - P\|_\infty \leq 10^{-15}$). The kernel operator was ported to `longdouble` and validated against the production `numba` operator to 3.8×10^{-15} , exactly the `float64` floor.

4.1 Headline: certified deficits across 140 cells

γ	$\tau=0.05$	0.10	0.20	0.30	0.40	0.50	0.60
0.05	5.9e-2	1.3e-1	1.7e-1	1.8e-1	1.8e-1	1.9e-1	1.9e-1
0.10	8.3e-3	5.4e-2	1.2e-1	1.5e-1	1.5e-1	1.5e-1	1.5e-1
0.27	1.3e-4	1.8e-3	1.0e-3	1.5e-2	5.4e-2	9.1e-2	9.1e-2
1.04	1.1e-6	1.7e-5	2.6e-4	1.2e-3	3.2e-3	6.7e-3	1.1e-2
5.57	6.3e-8	9.4e-7	1.3e-5	6.2e-5	2.1e-4	5.7e-4	1.3e-3
30.0	2.3e-9	3.2e-8	1.5e-6	2.5e-5	1.7e-4	6.7e-4	1.6e-3

Table 1: Certified revelation deficit $\mathcal{D} = 1 - R^2(\text{logit } P \sim \sum_k u_k)$ at selected (τ, γ) cells. All cells certified at $\|\Phi(P) - P\|_\infty \leq 10^{-15}$ in extended precision.




Figure 1: Certified deficit vs γ for each τ . Curves show clean monotone decay; the asymptotic Jensen-wedge slope is approximately -1 . The deficit hits the $G=21$ resolution floor below $\sim 10^{-7}$.




Figure 2: Same data on linear axes, zoomed to $\gamma \leq 3$. Shows the absolute size of the deficit where it matters economically.




Figure 3: Certified residuals $\|\Phi(P) - P\|_\infty$ in longdouble precision. All 140 cells are several orders of magnitude under the 10^{-15} certification bar.

5 Value-of-information decomposition

For each certified equilibrium we evaluate four primitives under the true joint distribution of (v, u_1, u_2, u_3) :

- $TV = \mathbb{E} |x^*(\mu_{\text{inf}}, P)|$: per-agent expected absolute trade.
- $\mu_{\text{FR}}(u) = \prod f_1(u_k) / (\prod f_0 + \prod f_1)$: full Bayes posterior given all signals (the welfare-relevant truth).
- $\mu_{\text{pub}}(p)$: rational price-only posterior, the Bayes update conditional on $\{P(U) = p\}$, computed by binning the joint distribution.
- $\mu_{\text{inf}}(u_k, p)$: the equilibrium's price-plus-own-signal posterior, again by binning.

Each baseline trades optimally at the equilibrium price under its posterior; certainty-equivalent wealth is evaluated under the true measure μ_{FR} . The decomposition is then

$$\begin{aligned} V_i^{\text{pub}} &= CE_{\mu_{\text{pub}}} - CE_{\text{prior}} \geq 0 \\ V_i^{\text{priv}} &= CE_{\text{informed}} - CE_{\mu_{\text{pub}}} \geq 0 \\ V_i^{\text{FRgap}} &= CE_{\text{FR}} - CE_{\text{informed}} \geq 0 \end{aligned}$$

with $V_i^{\text{pub}} + V_i^{\text{priv}} + V_i^{\text{FRgap}} = CE_{\text{FR}} - CE_{\text{prior}}$ the total CE-value of full information relative to the prior.




Figure 4: Value-of-information decomposition at $\tau = 0.20$. The price delivers the majority of the welfare value; the own private signal contributes a strictly positive but smaller share; the welfare gap to full revelation is tiny at high γ but $\Theta(V_i^{\text{pub}})$ at low γ .




Figure 5: Cell-level scatter: revelation deficit vs welfare gap V_i^{FRgap} . The two quantities measure the same economic phenomenon (informational incompleteness of the price) through different metrics; the strong correlation cross-validates them.

Key empirical findings.

1. $V_i^{\text{pub}} \gg V_i^{\text{priv}}$ at every parameterisation: the equilibrium price does the heavy informational lifting, the own private signal adds a smaller but *strictly positive* marginal benefit.
2. $V_i^{\text{priv}} > 0$ everywhere is the welfare counterpart of $\mathcal{D} > 0$: a positive deficit means the price is not a sufficient statistic for the signal triple, so private information is not redundant.
3. V_i^{FRgap} vanishes in proportion to $\mathcal{D}$ along each τ row, giving a direct welfare interpretation of the revelation deficit.

6 Asymptotic scaling and discussion

The Jensen wedge $\pi = \bar{m} + (\frac{1}{2} - p) \text{Var}(m)/\gamma$ between equilibrium price and posterior mean predicts $\mathcal{D} \propto 1/\gamma$ at large γ . Figure 1 confirms a clean asymptotic tail slope near -1 across all τ rows; finite-grid resolution and finite-precision noise dominate below $\sim 10^{-7}$. The wedge magnitude is increasing in τ (more informative signals produce more variance in posteriors and thus a stronger wedge), explaining why higher- τ rows have systematically larger deficits at every fixed γ .

Relation to the literature. Classical extensions of Grossman–Stiglitz [1, 2] introduce noisy supply to obtain partial revelation; our binary-payoff setting yields the same outcome through the Jensen wedge created by CRRA risk aversion, without any external noise. The closest precursors

work in CARA-normal economies; the CRRA-binary combination requires nonlinear clearing and creates the strict- $h=0$ degeneracy discussed in the appendix. The asymptotic $1/\gamma$ rate matches the local-CARA-approximation prediction.

7 Conclusion

Three CRRA traders observing Gaussian signals about a binary asset sustain a partially revealing equilibrium with no exogenous noise. We prove existence in a tilt class disjoint from full revelation and certify 140 equilibria at extended-precision residual $\leq 10^{-15}$, with per-cell trading volume and value-of-information decompositions consistent with theoretical predictions across six orders of magnitude in deficit.

A Strict- $h=0$ degeneracy under binary payoff

Setting the regularisation h to zero in the operator $\Phi_{h,\delta}$ exposes a degeneracy specific to the binary payoff. The fully revealing candidate $P_{\text{FR}}(u) = \sigma(\tau \sum_k u_k)$ is an exact fixed point of the strict- $h=0$ operator at every (τ, γ) : restricted to its own level set, the slice integrals close analytically, every trader’s posterior collapses to $\sigma(\tau \sum_k u_k)$, and CRRA clearing returns that same value. We verified this numerically in 9 independent cells spanning $\tau \in [0.05, 2]$ and $\gamma \in \{0.098, 1.035\}$, with the strict- $h=0$ residual at P_{FR} measured at $1.7 \times 10^{-15} \leq \|\Phi - P\| \leq 3.8 \times 10^{-15}$ (the float64 floor).

The implication is structural: the strict- $h=0$ fixed point problem admits the trivial revealing solution at *every* parameterisation and is in this sense degenerate. The partially revealing equilibrium of the main text exists only as a limit of fixed points of the regularised operator $\Phi_{h,\delta}$ as $h \rightarrow 0$. The kernel bandwidth h is thus *the equilibrium selection mechanism* for the nontrivial branch, not a numerical convenience. Proposition 2 shows the regularised operator never returns P_{FR} as a fixed point, so the selection mechanism is operative for all $h > 0$.

Section 4’s deep-ladder extrapolations along the joint-limit schedule converge to a limit cleanly distinct from P_{FR} : at $(\tau, \gamma) = (2, 0.098)$ the continuum estimate is $\mathcal{D}_\infty = 0.268 \pm 0.010$, far from 0. The economic content of the main text is therefore unaffected by the appendix’s degeneracy; the appendix simply records the technical reason why the limit construction is the principled route to the PR equilibrium under a binary payoff.

B Reproducibility

All data, scripts and certified solutions are in `projects/REZN/solved_fixed_points/lowtau/` (the 100-cell low- τ archive) and `projects/REZN/solved_fixed_points/dd_k3_overnight/emin15/` (the 40-cell mid- τ archive). Solver code: `projects/REZN/solver_code/highprec_dd_qd/`; metrics computation in `lowtau_metrics.py`; report builder in `lowtau_report.py`.

References

- [1] Grossman, S.J. and Stiglitz, J.E. (1980). On the impossibility of informationally efficient markets. *American Economic Review* 70(3): 393–408.
- [2] Vives, X. (2008). *Information and Learning in Markets*. Princeton University Press.