

Discretization Choices for the Strict- $h=0$ K=3 CRRA REE: Linear, $\sigma-\delta$ atanh, and CDF-cube Grids

Continuation of the methodology record — this session

May 31, 2026

Contents

1	Problem statement	1
2	Three cube discretizations	2
3	Two basins of the strict-$h=0$ operator	3
4	Existing reproduction: hfree_smooth at G=9	4
5	Extending to G=13	4
6	$\sigma-\delta$ V11 nails PR with Newton + FR-BC + warm-start	5
7	CDF-cube: a new discretization	5
8	Outlook: double-double precision	7
9	Summary of session findings	7

1 Problem statement

The noiseless K=3 CRRA rational-expectations equilibrium (REE) with binary asset $v \in \{0, 1\}$ at $\tau = 2$, $\gamma = 0.1$ admits a partially-revealing (PR) fixed point. The fundamental object is the equilibrium price $P(u_1, u_2, u_3) \in [0, 1]$ as a function of the three private signals. The fixed-point operator Φ is the Bayes-update + market-clearing map: agent k infers v from (u_k, P) , and the price clears the CRRA demands.

The contour-form evidence integrals are

$$A_v(p; u_k) = \int_{\{P(u_{-k})=p\}} f_v(u_{-k}) \frac{d\sigma}{|\nabla P|},$$

which require a discretization choice for the off-axis 2D slice $P(u_{-k})$. This paper documents three such choices for the strict- $h=0$ (no-kernel) operator class at this session's most-studied (γ, τ) .

2 Three cube discretizations

We compare three ways to lay out the 3D cube of unknowns P_{ijk} :

(1) Linear u -grid. $u_i = U_{\max}(2i/(G-1) - 1)$, uniform in $u \in [-U_{\max}, U_{\max}]$. This is the original `k3_hfree_smooth` representation. Cells are spaced linearly in u ; far tails are well-sampled but signal density $\bar{f}(u) = \frac{1}{2}f_0 + \frac{1}{2}f_1$ is concentrated near $u = \pm\frac{1}{2}$, so most of the linear grid sits in low-density regions.

(2) σ - δ atanh ξ -cube. Coordinates (u_1, Σ, δ) with $\Sigma = u_2 + u_3$, $\delta = u_2 - u_3$. Each axis is then atanh-stretched: $\xi_k = \tanh(u_k/T_k)$ for some half-width T_k , giving $u_k = T_k \tanh^{-1}(\xi_k)$ on the bounded cube $\xi \in (-1, 1)^3$. This is the σ - δ V11 (`v11_strict_h0_hardwired.py`) representation. The atanh squashing is independent of the signal distribution.

(3) CDF-cube (this proposal). Coordinates $\zeta_k = \bar{F}(u_k)$ where

$$\bar{F}(u) = \frac{1}{2} \Phi(\sqrt{\tau}(u + \frac{1}{2})) + \frac{1}{2} \Phi(\sqrt{\tau}(u - \frac{1}{2}))$$

is the unconditional signal CDF (Gaussian mixture). The cube $\zeta \in [\zeta_{\text{lo}}, \zeta_{\text{hi}}]^3$ is sampled uniformly in ζ , and $u_k = \bar{F}^{-1}(\zeta_k)$ is obtained by Brent root-finding on $\bar{F}$. By construction the grid nodes concentrate where signal mass is, giving a more faithful Riemann sum for the contour integrals than (1) or (2).




Figure 1: Three grid layouts at $G = 9$ shown in the (u_2, u_3) plane at fixed $u_1=0$, overlaid on the bivariate signal density (gray contours). Linear (blue) wastes nodes in the tails; σ - δ atanh (orange) is bounded but still mismatched; CDF-cube (green) tracks the density.




Figure 2: Left: the unconditional CDF $\bar{F}(u)$ and the $G = 9$ ζ -uniform nodes. Right: signal density with the three grids' node locations overlaid (stripes at $y = 0$ -0.05 linear, 0.05-0.10 atanh, 0.10-0.15 CDF). The CDF cube places nodes proportionally to $\bar{f}$; the atanh cube is too spread; the linear cube is too uniform.

3 Two basins of the strict- $h=0$ operator

A direct point of confusion at the start of this session was the apparent disagreement of the strict- $h=0$ operator with the (also strict- $h=0$) kernel co-area result. The resolution is that the $h=0$ operator admits *at least two* fixed points at this (γ, τ) :

- **PR basin:** slope ≈ 0.36 , deficit ≈ 0.17 , $d_{FR} \approx 0.26$. Reachable from the kernel co-area warm-start via Newton (`k3_hfree_smooth/hfree_nail.py`, commit 78c27216).
- **Near-FR basin:** slope $\in [0.85, 0.92]$, small d_{FR} , small deficit. Reachable from the no-learning IC via Picard.

Earlier in the session I had repeatedly tested the σ - δ V11 operator from no-learning Picard and observed the near-FR drift, mistakenly concluding that the strict- $h=0$ operator lacked a PR fixed point. The user's pointer to commit 78c27216 corrected this: with the right warm-start and a Newton-class solver, the PR FP is real, robust, and rock-solid.

Operator-class basin map at $\gamma = 0.1$, $\tau = 2$, $K=3$ CRRA
(GREEN: PR-basin success; RED: near-FR (wrong basin); ORANGE: degenerate/diverged)

Operator (h=0)	IC / warm-start	slope	d_FR	deficit	$\ F\ $	verdict
Kernel co-area h=0.32	no-learning	0.18	0.3	0.28	5e-11	PR (kernel-induced)
hfree smooth h=0 G=9	kernel-warm-start	0.36	0.26	0.17	4e-05	PR (TRUE h=0)
hfree smooth h=0 G=9	no-learning	0.92	0.05	0.04	?	near-FR (wrong basin)
hfree smooth h=0 G=1	G=9 interp (no Pic.)	0.41	0.23	0.13	0.19	PR (NK plateau)
hfree smooth h=0 G=1	G=9 interp + 5 Pic.	0.62	0.2	0.44	0.1	DEGENERATE
hfree smooth h=0 G=1	G=9 interp + Anderson	0.47	0.2	0.11	0.076	PR (best at G=13)
σ -6 V11 NK FR-BC G_in=1	hfree G=9 (interp)	0.55	0.18	—	8.4e-07	PR (nailed via NK)
σ -6 V11 NK warm-halo	hfree G=9 (interp)	nan	nan	—	0.045	DIVERGED
σ -6 V11 Picard G_in=1	no-learning	0.86	0.05	—	0.07	near-FR
CDF-cube Picard G=9	no-learning	0.5	0.13	0.12	0.07	PR (oscillates)

Figure 3: Operator-class basin map. Each row is an (operator, IC) pair; columns report final slope, d_{FR} , deficit and residual $\|F\|_\infty$ achieved. Green = PR-basin success; red = wrong (near-FR) basin; orange = divergent or degenerate halfway state. The same operator nails or fails depending on the initial-condition basin and solver.

4 Existing reproduction: hfree_smooth at G=9

Re-running the original Newton-Krylov on `P_nailed.G9.npy` this session confirms the published result of commit 78c27216: 3 NK iterations, $\|F\|_\infty$ trajectory $8.4 \times 10^{-7} \rightarrow 1.2 \times 10^{-8} \rightarrow 3.1 \times 10^{-11}$, final deficit 0.1725, $\text{slope}_T = 0.3641$, $d_{FR} = 0.2566$. This is the canonical strict- $h=0$ PR fixed point at $G = 9$ on the linear u -grid.

5 Extending to G=13

Continuing the same operator to $G = 13$ proved hard. The PR basin at $G = 13$ is narrower and the standard “interpolate + Picard precondition + NK” recipe often kicks the iterate out of basin. We tried three warm-start strategies:

- **A1 (raw interp, 0 Picard, then NK):** stays in PR basin (slope 0.41, $d_{FR}=0.23$) but NK plateaus at $\|F\| = 0.19$ — LGMRES saturation.
- **A2 (5-iter Picard pre-conditioning, then NK):** jumps slope $0.43 \rightarrow 0.76$ in one NK step and lands in a degenerate halfway state (deficit 0.44). Even a few Picard iterations were enough to escape PR.

- **A3 (Anderson m=8 Picard from raw interp):** best result. Reaches $\|F\| = 0.076$ at iteration 6, slope 0.47, deficit 0.11. Stays in the PR basin but doesn't fully nail.

The $G=13$ result is the best PR-basin solution obtained on the linear u -grid at $G = 13$ to date, but the residual $\|F\| = 0.076$ is far from the discretization floor. Either a smarter continuation in G (pseudo-arclength) or higher precision (see Sec. 8) would be needed for a tight nail.

6 σ - δ V11 nails PR with Newton + FR-BC + warm-start

After ruling out boundary-condition causes (Sec. 3 of the earlier `extended_paper.tex`), this session also re-tested σ - δ V11 on the right warm-start. The recipe is:

1. Trilinear-interpolate the `hfree_smooth` $G=9$ PR FP onto the σ - δ ξ -cube ($G_{FULL}=13$, $G_{inner}=11$).
2. Apply FR boundary conditions at the $\xi = \pm 1$ faces ($P = 0$ or 1 exactly — physically the perfect-info limit at $u_k = \pm\infty$).
3. Run `scipy newton_krylov` on the inner-block residual with $f_tol=10^{-7}$.

Result: NK iter 1 reaches $\|F\| = 8.4 \times 10^{-7}$, below tolerance. The same operator that earlier “drifted to near-FR” under no-learning Picard *nails the PR FP in a single NK iteration* given the right warm-start.

A control variant in which the halo is held at the interpolated warm-start (instead of FR limits) diverges to NaN within a few NK iterations: the FR limit at $\xi = \pm 1$ is the unique consistent boundary condition for the strict- $h=0$ operator on the ξ -cube.

7 CDF-cube: a new discretization

Section 2 introduced the CDF-cube as a third discretization. The motivation is that the strict- $h=0$ operator depends on contour integrals dominated by the signal density's support; sampling proportionally to that density should reduce quadrature error and (empirically) widen the PR basin's stability region.

Initial test at $G = 9$, $\zeta \in [0.001, 0.999]$, $u \in [-2.54, 2.54]$:

- No-learning IC: deficit = 0.40, slope = 0.13, $d_{FR} = 0.31$ — already inside the PR basin (in contrast to the linear u -grid IC which sits in the near-FR basin under Picard).
- Plain Picard, iter 1: slope $0.13 \rightarrow 0.45$, d_{FR} $0.30 \rightarrow 0.22$. For comparison, linear u -grid Picard at the same step moves slope $0.19 \rightarrow 0.73$ and σ - δ V11 moves $0.25 \rightarrow 0.79$ — both escape PR in one step, the CDF-cube stays.
- Picard iterations 1-11 oscillate gently around slope ≈ 0.48 , $d_{FR} \approx 0.13$, $\|F\|$ between 0.07 and 0.50 (alternating). The operator clearly has a PR FP on this grid; plain Picard isn't tight enough to nail it.

Switching to Newton-Krylov with continuation in γ is in progress; the first $\gamma = 0.01$ tile already shows a PR-like NL IC. (Sweep results will be incorporated in a follow-up revision.)




Figure 4: Residual $\|F\|_\infty$ trajectories across the operators and warm-start strategies considered this session. The dense Newton at $G = 9$ (blue line) is the gold-standard quadratic convergence; NK re-verifies it. $G = 13$ A3 Anderson (purple) is the best PR-basin $G=13$ nail; A1 (green) plateaus; A2 (orange) escapes basin. σ - δ NK FR-BC (red) nails the PR FP in one iteration from the right warm-start.




Figure 5: The PR fixed-point $P(u_2, u_3)$ slice at $u_1=0$ for (1) hfree linear $G = 9$ (slope $_T=0.36$), (2) kernel co-area $h \approx 0.32$ $G = 17$ (slope $_T=0.18$). The CDF-cube panel will fill once the γ -sweep first save lands. Color: blue $P \approx 0$, white $P \approx 0.5$, red $P \approx 1$.

8 Outlook: double-double precision

At the smaller grids ($G = 9$) the strict- $h=0$ PR FP nails to $\|F\| \sim 10^{-11}$ with float64 + Newton. At $G = 13$ and above, two precision issues bite:

- The PR basin narrows. Small numerical errors in the operator evaluation become comparable to inter-basin distances, and Newton steps may exit the basin during line search.
- LGMRES inner solver hits a precision-induced floor (here ~ 0.02 at $G = 13$) below which the iterate cycles in numerical noise.

A double-double (or higher) operator implementation — already developed in `projects/REZN/solver_code/hfree` for the $K=2$ problem — should both narrow the precision floor and stabilize the PR basin at large G . Porting the CDF-cube operator (Sec. 2) to DD is a natural next step if the float64 NK γ -sweep saturates.

9 Summary of session findings

1. The strict- $h=0$ $K=3$ CRRA REE *does* have a PR fixed point at $(\gamma, \tau)=(0.1, 2)$ on the linear u -grid at $G = 9$, with deficit 0.17, slope 0.36, $d_{FR} = 0.26$. Confirmed by NK re-verify from existing `P_nailed_G9.npy` to $\|F\| = 3 \times 10^{-11}$.
2. The same operator's near-FR basin (slope ≈ 0.86) is what no-learning Picard reaches; this is what mislead earlier σ - δ V11 / u -grid / `crxa_xi_pchip` tests. *Initial condition matters as much as operator class.*
3. σ - δ V11 + Newton-Krylov + FR boundary + `hfree` warm-start nails the PR FP to $\|F\| = 8.4 \times 10^{-7}$ in one NK iteration — proving the (u_1, Σ, δ) frame admits the same PR FP.
4. Extending `hfree_smooth` to $G = 13$ is sensitive to warm-start. Anderson Picard stays in basin best ($\|F\| = 0.076$, slope 0.47); a single Picard preconditioning followed by NK escapes basin.
5. The CDF-cube proposal (this work) yields a wider PR basin: NL IC is already PR-like, and Picard moves slope $0.13 \rightarrow 0.45$ in one step (vs $0.19 \rightarrow 0.73$ for linear u -grid). NK + γ -sweep in progress.

Reproduction

All code and intermediate artifacts on branch `claude/study-fixed-point-economics-y12PB`. Key scripts:

- `projects/REZN/solver_code/sigma_delta/cdf_cube_operator.py` — the CDF-cube CRRA Phi.
- `projects/REZN/solver_code/sigma_delta/cdf_cube_NK_sweep.py` — NK + γ continuation.
- `projects/REZN/solver_code/sigma_delta/sigdelta_nk_from_hfree.py` — σ - δ V11 NK from `hfree` warm-start.
- `projects/REZN/solved_fixed_points/k3_hfree_smooth/g13_nail_clean.py` — $G=13$ three-attempt nail.